

Intro to Multi-Threading

Sam Shue, PhD



Demo Code and Slides:


<https://github.com/SamShue/Intro-to-MultiThreading>



THE WILLIAM STATES LEE
COLLEGE OF ENGINEERING

Demo Code and Slides:

<https://github.com/SamShue/Intro-to-MultiThreading>




I have an application in C++ that extracts edges from several user-specified images.



Gaussian Blur
+
Laplacian
Edge
Detection



Current Application Flowchart






Application Demo

Problem: The process of detecting the edges takes a long time, and so does the user typing in the file name.

Desire: I would like to collect user input while the program also processes the image in the background

Desired Application Flowchart




How do we get our application to simultaneously execute different lines of code?

Answer: **Threads!**

Threads allow lines of code to be executed simultaneously

We can organize our code into functions, give the thread the function to execute, and continue executing other code while the thread works at the same time



How do we make a thread in C++?

Step 0: Include standard thread library

```
#include <thread>
```

Step 1: Organize simultaneous behavior into functions

```
// A dummy function  
void foo(int Z)
```

Step 2: Create thread object and give it the function to execute


```
thread th1(foo, 3);
```



What do we do if we need data that a thread is operating on?

We can wait for a thread to finish with the “join” function

The “join” function will stall the code until the thread is finished



```
th1.join();
```



C++ Thread Demo

Let's thread our image processing application

We'll functionalize the edge detection algorithm and file access

A thread will then call the function once the user specifies the file



Threaded Application Demo

Success!


The user input can be collected
while at the same time the image is
processed!

But wait... what happened to our
images?


Our image processing functions share a temporary buffer when working on images...

When multiple threads run at the same time, they both use this same buffer and write over each others' data!

This is called a race condition



Solution: We need to lock the data!



This kind of lock is called a mutex
(short for **MUT**ual **EX**clusion)

A mutex can be locked by a thread
to secure a resource

If another thread tries to lock a
mutex that is already locked, it will
stall until it can acquire the lock



How do we make a mutex in C++?

Step 0: Include mutex library

```
#include <mutex>
```

Step 1: Create global mutex object

```
mutex m;
```

Step 2: Lock and unlock the mutex
to secure the resource

```
m.lock();
```

```
m.unlock();
```



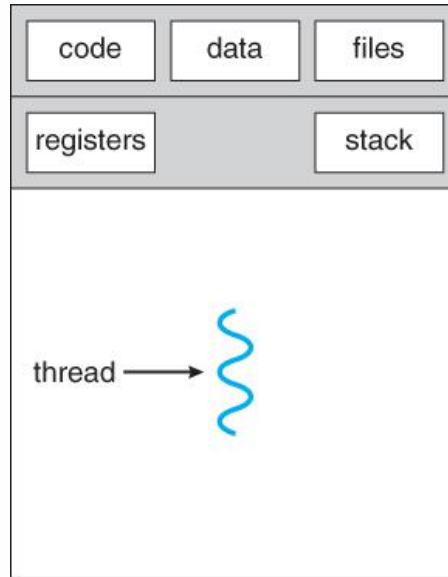
Mutex Locked and Threaded Application Demo

How does the computer simultaneously execute multiple threads?

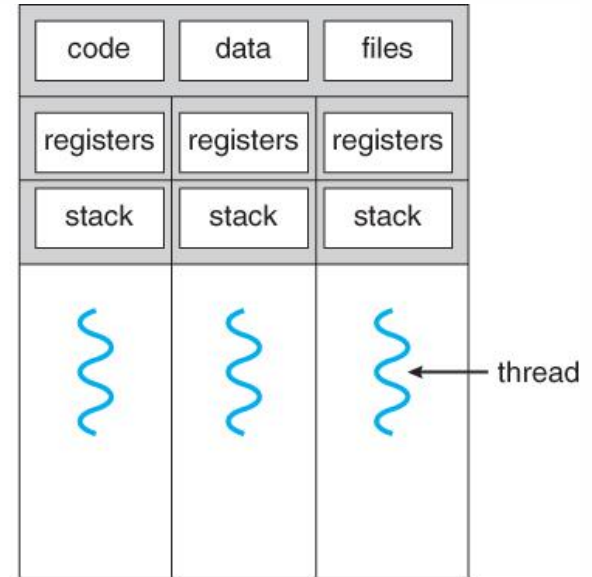
The **operating system** manages memory for each thread and then swaps back and forth between each thread

The operating system creates separate stack space for each thread

Other memory space is shared (heap, data, bss)




single-threaded process



multithreaded process

The operating system's thread scheduler lets each thread have execution time on the CPU





A single-core CPU
will have to run
threads concurrently


Multi-core CPUs can
run threads in
parallel

Parallelism

Thread 1 - Core 1



Thread 2 - Core 2



Concurrency

Thread 1 - Core 1



Thread 2 - Core 2



blocked on critical section
(Mutual exclusion)


Executing
critical
Section



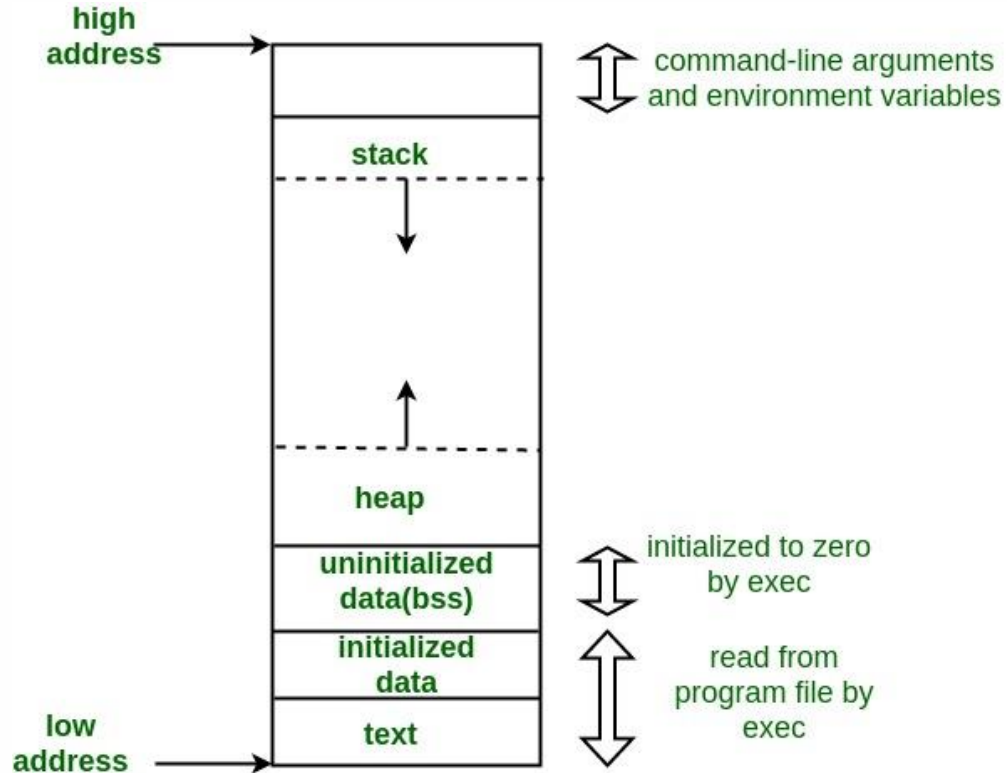
Thank You for Coming!

Questions?

Appendix A:




Appendix B:






Appendix C:




Appendix D:

Without hyperthreading



With hyperthreading





Appendix E:

MUTEX



SEMAPHORE



four identical keys

